

arraylist example

```
int x;
arraylist_t list; // local variable, but will refer to heap data
al_init(&list, 10);
// we passed a pointer, so al_init can change the actual fields
// of this local variable

al_remove(&list, &x); // remove last item, write value to x

al_remove(&list, NULL); // remove last item, discard value
```

Why do we pass the arraylist struct by reference?
i.e., why do we pass a pointer?

-> because we want our changes to persist after the function returns
-> passing a copy of the struct would mean that we can't change used

What is the operating system for?

-> provides an interface between programs and the hardware
-> abstracts details of the hardware
-> isolates programs and users
-> multiple programs can run without interfering with each other

therefore, program IO almost always happens via OS services ("system calls" "syscalls")

-> instead of taking directly to the disk or display hardware,
we ask the OS to send byte sequences appropriately

How system calls happen is different depending on the HW and OS

-> C provides functions that wrap the syscall
-> the functions themselves are HW/OS specific and may be written in assembly

Posix standardizes system calls for Unix-like operating systems

there used to many, many proprietary Unix variants
having a standard made it easier to write portable code
-> nowadays, Unix is pretty much always Linux or BSD (including MacOS)
-> plus a bunch of other software (eg. GNU utilities)

The Posix functions for file IO are open, close, read, write, &c.

-> central interface is the file descriptor
-> this is a number that indicates which file you are talking about

Each process has a table of open files; the file descriptor is an index into this table

-> remember: files can be any stream of bytes
-> e.g., we could read from a file on disk, or a terminal interface, or the network
or other processes

```
#include <fcntl.h> // "file control"
provides open and close
```

```
int open(char *filename, int flags)
```

filename is the name of a file or a path to the file
flags is a bit-vector that says what features we want

Must start with one of these
O_RDONLY - open a file in read mode
O_WRONLY - open a file in write mode
O_RDWR - open a file in read/write mode

Use bit-wise or to add additional features

O_APPEND - start with file pointer at the end of a file
O_TRUNC - delete contents of file
O_CREAT - create file if it does not exist (*special)

to open a file in append mode, we would say

```
int fd = open("log.txt", O_WRONLY|O_APPEND);
```

Use "man 2 open" to see a list of possible flags

most of time we will open in read mode

```
int fd = open("my_input.txt", O_RDONLY);
```

open returns -1 if it could not open the file
file descriptors are always non-negative
when open fails, it will set the global variable errno

traditionally, file descriptors 0, 1, and 2 are already open when your program starts

0 is standard input
1 is standard output
2 is standard error
you don't open these yourself

Using O_CREAT adds a third parameter to open, so we can specify its permissions

```
int open(char *filename, int flags, mode_t mode)
mode indicates what permissions the file should have
e.g., user/group/world readable/writeable/executable
we will come back to this
```

Use close to close a file

```
int fd = open("my_file", O_RDONLY);
...
close(fd);
```

Reading from a file

```
#include <unistd.h>
```

```
size_t read(int file_descriptor, void *buffer, size_t buffer_size)
```

-> reads from the specified file
-> writes bytes starting at the address specified in buffer
-> reads up to buffer_size bytes
-> returns the number of bytes read
-> returns 0 at end-of-file
-> returns -1 and sets errno if something went wrong

```
// read from stdin
```

```
int bytes_read;
char buf[256];
```

```
bytes_read = read(0, buf, 256);
// read UP TO 256 bytes from standard input (0)
// we might read fewer than 256 bytes, but we definitely won't read more
// write bytes to buf (an array of chars)
// return the number of bytes successfully read (assigned to "bytes_read")
// if we ask for more bytes than there are available, we get all of them
// if only 100 bytes are available, then bytes_read will be 100
```

read does not read strings! If we want a terminator, we have to add it ourself

```
// writing to a file
```

```
ssize_t write(int file_descriptor, void *buffer, size_t count)
-> writes to the specified file
-> bytes are taken starting at address specified (buffer)
-> number of bytes to write is specified (count)
-> returns the number of bytes successfully written
-> returns -1 and sets errno if something went wrong
-> note that the buffer can be any type!
```

```
// write to stdout
```

```
int bytes_written;
char buf[256] = "Hello!\n";
```

```
bytes_written = write(1, buf, strlen(buf));
// write to stdout (file descriptor 1)
// take bytes from buf
// only write the bytes in the string, not including terminator
// bytes_written will be assigned how much we were able to write
```

Q. why didn't we do this?

```
write(1, buf, 256);
-> this would write all the garbage data after the \0

write(1, buf, strlen(buf) + 1);
-> this would write the terminator
-> ... but we don't normally want to do this
```

Remember: terminators are used in C strings, but files are not strings
files can contain many \0 bytes, or none
file length is tracked by file system

Using read and write with binary files

-> read and write work with any data
-> we take bytes directly from memory without interpretation

Writing an array of ints to a file

```
int a[] = {1,2,3,4};
```

```
int fd = open("my_data", O_WRONLY|O_TRUNC|O_CREAT, S_IRUSR);
```

```
write(fd, a, sizeof(int) * 4);
// writes the array of ints to the file directly
// does not print as text! does no interpretation!
// the literal bytes from memory are written to the file
```

the file will be (say) 16 bytes,

on a little-endian machine (Intel) with 4-byte ints, it will contain (in hex):

```
01 00 00 00 02 00 00 00 03 00 00 00 04 00 00 00
```

on a big-endian machine (ARM) with 4-byte ints, it will contain (in hex):

```
00 00 00 01 00 00 00 02 00 00 00 03 00 00 00 04
```

this is why file transfer between different hardware can be challenging!

We can use read in the same way

```
int a[4];
int fd = open("my_data", O_RDONLY);
read(fd, a, 4 * sizeof(int));
```

```
// reads 16 bytes from the file and writes the bytes to memory in a
// no interpretation!
```

Typical pattern when reading a text file

two loops
outer loop calls read and writes to a char buffer
inner loop iterates through the bytes in the buffer
outer loop ends when read returns 0 (EOF) or -1 (error)

```
// int bytes_read, length;
// char *buffer or char buffer[...]
```

```
bytes_read = read(my_file_desc, buffer, length);
while (bytes_read > 0) {
    for (i = 0; i < bytes_read; ++i) {
        // do something with buffer[i]
    }
    bytes_read = read(my_file_desc, buffer, length);
}
```

```
// note that we have to keep reading until we get 0
// don't want to assume how big the file is
// allocating a giant buffer is wasteful
```

Some people prefer not to duplicate the call to read

```
while ((bytes_read = read(my_file_desc, buffer, length)) > 0) {
    for (i = 0; i < bytes_read; ++i) {
        // do something with buffer[i]
    }
}
```

Recall that assignment is an expression that evaluates to the value assigned

```
(x = foo()) == y
// calls foo(), assigns return value to x, and compares with y
```

Some people feel that putting the call to read in the while loop condition is
confusing and makes it hard to see what is happening

Reading:

man pages for open, close, read, write
these are in section 2 of the manual, so use "man 2 read", etc. to make sure
you get the right one