

=====

Final exam:

May 10, 4:00 PM - May 11 4:00 PM
 approx. 2 hours (may change)
 offered through Sakai
 primarily multiple choice
 -> "don't know" gives you 25% credit

Review of topics

- data types

- determining size of data: sizeof()
 sizeof(int) sizeof(struct addrinfo)
 sizeof(int *) sizeof(struct addrinfo *)

 sizeof(variable) <- don't be confused!

 int array[20];
 sizeof(array) == 20 * sizeof(int)

 int *array = malloc(20 * sizeof(int));
 sizeof(array) == sizeof(int *)

- strings vs char arrays

- strings always end with a terminator ('\0')
- character arrays may have no terminator or many terminators
- the difference is how they are used

- when we call read() we get raw data; not necessarily a string

- strcpy(), strcmp(), strdup(), strlen(), strcat()
 strncpy(), strncmp(), strndup(), strnlen(), strncat()

- "n" variants take a maximum length
- useful if we can't guarantee the presence of a terminator
- useful if we aren't sure destination of strcpy() is big enough

- creating types: struct, union, enum, typedef
 struct bundles together multiple values
 union gives us a choice between types - hard to use correctly

- enum - fancy integer constants
 enum cardinal_direction { north, south, east, west };

- typedef - abbreviate type names
 typedef void *(*thread_fun)(void *);

- now we can write
 int pthread_create(..., thread_fun, ...);

- pointers

- size of object being pointed to
- assigning to pointer variable vs copying data

```
int *p;  
p = q; // makes p point to the same int as q
```

```
*p = *q; // copies the value q points to into the place p points to
```

```
char *msg = "Hello";  
msg = "Goodbye"; // changes which string literal msg points to  
strcpy(msg, "Goodbye"); // segmentation violation!  
// msg points to a string in the data segment (read-only memory!)
```

```
char buf[] = "Empty"; // buf names an array of 6 chars  
buf = "Full"; // not allowed!  
strcpy(buf, "Full"); // copies characters to buf
```

```
char *nothing;  
strcpy(nothing, "Something"); // wrong!  
nothing does not point to anything
```

```
char *something = malloc(10);  
strcpy(something, "Something"); // ok, because we allocated enough space
```

```
char *somethingelse = malloc(10);  
somethingelse = "Something"; // memory leak!  
// we no longer have a reference to the 10-byte object from malloc()  
// therefore, we have no way to free() it
```

- dereferencing (*)
 left side: *p = x; // writes to the object p points to
 right side x = *p; // writes the value that p points to

- address-of (&)
 &variable - address of a variable
 &array[index] - address of specific item in array

```
array[0] - value of first element  
array - address of array / address of first element  
&array - address of array / address of first element  
&array[0] - address of first element
```

```
p[n] == *(p + n)  
- prefer array index notation to pointer arithmetic
```

```
&struct_ptr->field - address of field in structure pointed to by struct_ptr
```

- arithmetic: incrementing/decrementing pointers

```
some_type *p = malloc(sizeof(some_type) * n);
```

```
p - address of first item / object that p points to  
p + 1 - address of value after the one that p points to
```

```
*p - value of object that p points to / first item in array  
*(p + 1) - value of object after the one that p points to  
- same as p[1]  
- also the same as 1[p], if you like confusing people
```

- contrast with array variables
 array variables are always static or on the stack

```
pointer can point to anything  
  
array variables cannot be reassigned  
pointers may be freely reassigned
```

```
compiler knows how big the array for an array variable is  
only you know how big an array a pointer points to
```

- void *
 point to anything; effectively the same as char *

```
C automatically casts pointers to and from void *  
we can't meaningfully dereference them  
- we cast to the right kind of pointer and then use that  
- compiler assumes that what we are doing is correct  
- no guaranteed behavior if we accidentally cast to the wrong type
```

- variables

- scope/visibility
 - associated with a particular block or the top-level

- visible after being declared
- variables in an inner scope "shadow" variables with the same name in outer scopes

```
int foo; // global  
  
void fun(void)  
{  
    int foo; // local; shadows global foo (no way to refer to global foo)  
  
    foo = 5; // refers to local foo, not global foo  
}
```

```
void fun2(void)  
{  
    foo = 5; // refers to global foo  
}
```

- binding with object
 top-level (global) variables bind to static (process-lifetime) objects
 local variables are usually stack (function-lifetime) objects
 static local variables have process lifetime but local scope

- initialization vs assignment

```
int a[4] = {1,2,3,4}; // array initializer  
a = {1,2,3,4}; // syntax error (cannot use initializer syntax for assignment)
```

```
int *p = "Foo"; // pointer initializer (p points to string literal)  
p = "Foo"; // pointer assignment (p points to string literal)  
*p = "Foo"; // type error (*p expects char, but "Foo" is char *)
```

- objects: anything stored in memory

- lifetime: static, function, arbitrary
 static - stored in data segment; exists while process runs

- function - stored in stack; created when function starts, destroyed when function exits

- arbitrary - stored in heap; created by malloc(), destroyed by free()

- location: data segment, stack, heap
- direct reference (variable name) vs indirect reference (pointer)

- malloc(), calloc(), realloc(), free()
 - avoiding leaks
 - detecting errors

```
p = realloc(p, new_size); // risky! if realloc() fails, we lose our pointer
```

- memcpy(), memmove(), memset()

- file IO

- buffered operations (FILE *, fopen(), fprintf(), etc.)
- non-buffered operations (int, open(), write(), etc.)

- modes: read, write, append, read/write
- file pointer/cursor: keeps track of where we last read/wrote

- control with lseek(), fseek()
- use of file IO for non-files: pipe(), socket(), accept()

- some operations restricted to specific kinds of file
- opendir(), readdir(), closedir(), struct dirent

- stat(), fstat()
- dup(), dup2() - additional file descriptor for existing open file

- preprocessing, compilation, and linking

- what goes in a header file
- #include vs linking
- #define macros

- processes

- process ID
- fork() - create a child process
- wait() - wait for a child process to terminate

- zombie process, orphan process, zombie-orphan process
- execl(), execv() - change the program a process is executing

- many things shared between parent and child process
- what happens if we fork() when a file is open?

- many process attributes are preserved when we exec()

- signals

- setting a signal handler
- signal(), sigaction()
- blocking signals

- threads

- pthread_create(), pthread_join(), pthread_detach()
- if nothing joins or detaches a thread, it becomes a zombie when it stops

- coordination
- mutex, condition variable, barrier, semaphore

- deadlock: 4 necessary conditions
- mutual exclusion
- hold and wait
- no preemption
- circular wait

- file system

- inodes

- use of indirection to allow:
 1. constant size of inode
 2. minimal wasted space for small files
 3. possibility to represent very large files

- directories (association of names to inodes)
- paths
- paths-inode relation is many-to-one!
- file permissions/modes

- The command line

- running programs: bare name vs path
 \$ program_name
 use search path to find program
 \$ path/to/program_name
 \$./program_in_current_dir
 specifies program; no need to search

- file globs
- controlling stdin/stdout/stderr: pipes, redirecting

```
cat file | ./ww 20  
output from first program sent as input to next program  
./program < input  
./program > output
```

- utility programs

- man
- cat, more, less, head, tail
- file, wc
- cmp, diff
- grep
- ps, top, jobs
- tee
- echo

- file management utility programs/commands
- ls, cp, mv, rm
- cd
- mkdir, rmdir
- chmod, chown

- networking

- 7-layer OSI model
- 4-layer internet model
- addresses at various layers

- link layer: identify specific machine on local network (e.g., MAC)
- network layer: identify machine globally (IP address)
- transport layer: identify process on machine (IP address + port)
- application layer: identify things relevant to application (e-mail address, URL)

- not all devices interact with the entire stack
- switches and hubs live at link level; invisible to network and above
- routers live at network layer; mostly invisible to transport and above
- user programs live in application layer (sockets talk to transport layer but do not expose details)

- sockets

- listening vs connection sockets
- listening - can use accept()
- connection - can use read(), write()

- socket(), bind(), listen(), accept()
- socket(), connect()
- domain name vs IP address; getaddrinfo(), getnameinfo()
- what can go wrong?

Remember to review the course and your recitation on SIRS!