

Regarding dirent.h and readdir

The only field that Posix requires in struct dirent is d\_name  
d\_name is guaranteed to point to a terminated string  
NAME\_MAX from limits.h is the maximum length of this string  
It is not recommended to create an array of length NAME\_MAX  
NAME\_MAX may be extremely long, depending on your OS/architecture

If you are using -std=c90 or -std=c99, you won't have access to other fields, like d\_type  
-> but you can use stat to get this information

If you are using GNU extensions (the default for GCC), then d\_type is available  
-> it does not matter whether you use d\_type or stat

d\_type is technically non-standard, but it is available on BSD and Linux systems, which is most of the Unix-like systems you are likely to encounter  
  
d\_type is faster than stat, because you don't need to call the OS again but the speed difference is not significant for us

---

How are you verifying your program?  
- just eyeballing the results  
- made a bunch of sample output and compare program output

- write a program to test the output of your program  
- what properties should your program output have?  
- Only three whitespace sequences: " ", "\n", "\n\n"  
- in particular, no spaces adjacent to newlines  
- The number of characters between two newlines should not exceed the page width  
- exception if input contains words that are too long  
- If you run ww again on the output of ww with the same width, you should get the same text  
  
./ww 80 some\_input > output1  
./ww 80 output1 > output2  
  
output1 and output2 should be identical

Useful unix command: tee

\$ tee <file\_name>  
  
tee reads from standard input and writes to standard output and a file  
lets me split the input and send two copies to different places  
handy for saving data from the middle of a pipeline  
  
\$ program1 | tee p1\_out | program2 > p2\_out  
  
Output from program1 is sent to tee  
tee passes the output to program2 and saves a copy to p1\_out  
Output of program2 is written to p2\_out

---

Reminder about paths

-> the file names we give to open, opendir, fopen, etc. are really paths

A path is a sequence of names separated by slashes (/)  
Paths starting with / are absolute  
Other paths are relative  
A relative path says how to reach a file, starting from the current directory (also called the working directory)  
  
Eg., foo/bar/baz

means "baz" inside "bar" inside "foo" in the current directory  
  
Our processes have a current working directory:  
the directory the user was in when they started the program  
that is, we inherit the working directory from the shell  
  
-> thus, file names are interpreted relative to the directory we started in

How can we open files in another directory?

1. make a path: directory/filename  
when reading from a directory, just append a slash and the file name to the name we used to open the directory  
-> a little obnoxious, since strcat requires us to allocate space in advance  
-> but eminently doable  
  
There is a limit on path length, but this is not generally a problem  
-> open will just return -1 if the name is too long  
-> the max path length may be very long, so use dynamic (heap) allocation  
  
2. use chdir to change the working directory  
now all my filenames are relative to the new directory  
no need for concatenating path names  
  
int chdir(char \*path);  
  
returns 0 for success  
on error, returns -1 and sets errno

--

/dev directory contains "files" that correspond to hardware devices

/dev/null - a device that just throws away its input  
- commonly used as a way to throw out/ignore the output from a program  
  
./program\_with\_lots\_of\_output > /dev/null

On Linux, /proc contains an entry for every running process

---

Processes  
-----

What is a "process"?  
- for us, a process is a program that is currently executing  
  
process = program + state  
  
state includes contents of memory, registers, open files, OS data

In principle, we can have multiple processes running concurrently that all use the same program

The OS keeps track of each running process  
On Unix, we have the Process Control Block  
-> data structure used by the OS  
- process ID (pid)  
- what virtual memory address space we are using  
- what files we have open  
- environment information (e.g., current working directory)  
- permissions, process user/group, etc.

Where do new processes come from?  
In general, the OS has some mechanism  
In Unix-like systems, new processes are created by fork()

pid\_t fork();  
  
pid\_t is an integer type (width is OS-dependent)  
  
fork clones or duplicates the current process  
-> the original process is called the "parent"  
-> the new process is called the "child"  
  
The OS has a special program that starts first (possibly "init")  
every other process is the child of some other process  
  
When we call fork, the current process is duplicated  
- both duplicates have the same program  
- initially have the same memory contents / program state  
- but the child gets a copy of the memory  
- after fork returns, the parent and child may diverge  
  
What fork returns will be different between the parent and child  
-> parent receives the PID of the child  
-> child receives 0  
(0 is a valid PID, but it always identifies the root process, so it can never be the PID of a child process)  
  
Typical usage  
  
pid\_t pid = fork();  
if (pid < 0) {  
 perror("fork");  
 abort();  
} else if (pid == 0) {  
 puts("In child!");  
} else {  
 printf("In parent! Child is %d\n", pid);  
}  
  
Remember: both child and parent start with the same program in the same state (aside from the return value from fork)  
we can think of fork as returning twice  
-> once to parent  
-> once to child  
both parent and child resume executing after the return from fork  
  
Fork can fail if the OS (or the user) has too many processes  
fork returns -1 and sets errno on failure

What happens when a process ends?  
-> who collects our return code?  
  
- Much of the data for our process is forgotten  
- The process control block holds our exit status, in case our parent wants it

When a process ends, it becomes a "zombie"  
A zombie process has died, but it's PCB still lives in the process list

The parent must call wait() to clean up the zombie process  
wait is also how we get the exit status of the child

If a parent process ends before its child, the child becomes an "orphan"  
An orphan process has no parent process to wait for it when it ends  
-> OS will make sure the orphan is "adopted" by some process (usually init)  
- the adopter will wait for the child to terminate

If the parent and child both terminate, the child becomes a "zombie orphan"  
-> The OS will have some process adopt the child and wait

Only specialized processes can adopt orphans  
- process do not necessarily adopt distant descendants  
- there is no "grandparent" relationship

wait  
----  
  
pid\_t wait(int \*wstatus);  
  
wait suspends a program (blocks) until one of its child processes terminates  
It returns the PID of the terminated process  
If wstatus is not NULL, it writes information about how the process exited into wstatus  
-> exit status  
-> whether the process was terminated by a signal  
  
wait (and related functions) are how we prevent orphans  
  
any time we call fork, the parent should call wait  
  
When should we call wait?  
It depends on what we're doing

Why does any of this matter?  
Next time: exec and multiprocessing