

Recap: fork and exec

- fork "returns twice" in two different processes
  - fork creates a duplicate of the current process
    - the duplicate is called a child, the original is the parent
    - both start with the same contents of memory, program counter, etc.
    - as though both processes had been running since the start and behaving identically
  - in the child, fork returns 0
  - in the parent, fork returns the PID of the child
- the parent must use wait to clean up the child's PCB
  - this also tells the parent how the child exited
    - what was the exit status
    - was the child terminated by a signal
- execl and execv change what program a process is executing
  - specify file name for program we want to start executing
  - provide arguments explicitly
  - much of our other process information is preserved
    - e.g., open files
- exec \*changes\* the current program
  - what we are currently doing is not preserved
  - exec does not return (unless it failed)

we can take advantage of this to do things like

- spawn a process and read what it writes to standard output
- use pipe to create a pair of file descriptors (read end and write end)
- use fork to start a child process
  - the pipe is now shared between the parent and child
- in the child
  - use dup2 to set the write end of the pipe to be standard output
  - use execl or execv to start our new program
  - > anything the new program writes will get sent to the pipe
- in the parent
  - read from the pipe to see the child program's output
  - wait for the child once we are done reading
- for two-way communication, use two pipes

recap: wait

- wait pauses the process until a child process ends
  - > if a child process has already ended, it returns immediately
  - > if there are no child processes, it returns an error
  - > if there is more than one child, it waits for the first one to end

related:

- timedwait if you don't want to wait more than some amount of time
- waitpid if you want to wait for a specific child

```
{
    pid_t child1 = fork();
    if (child1 == 0) { ... execv(program_1, plargs); ... }

    pid_t child2 = fork();
    if (child2 == 0) { ... execv(program_2, p2args); ...}

    // at this point, we have two child processes (assuming no errors from fork)
    // we have to call wait twice to clean up both children

    // the two children run concurrently with us

    pid_t finished[2];

    finished[0] = wait(NULL); // wait for one child to exit
    finished[1] = wait(NULL); // wait for other child to exit

    // finished tells us which child ended first
    if (finished[0] == child1) {
        puts("Child 1 was faster");
    }
}

pid_t wait(int *wstatus);

    return value -> PID of next child to terminate (or -1)
    exit information will be written to wstatus
    use macros like WEXITSTATUS to get information from wstatus
```

reminder: fork returns the PID of the child process  
to get the exit status of the child, we must use wait

```
if (fork() == 0) {    // FIXME we should also check for -1
    // do child stuff (e.g., exec)
    exit(1); // just to make sure we terminated (shouldn't be necessary)
}

int wstatus;
pid_t finished = wait(&wstatus); // wait for child to finish
// FIXME: should be checking for error return value

// finished contains PID of child
// WEXITSTATUS(wstatus) is the exit status of the child
```

What happens if we call fork multiple times?

```
fork();
fork();
fork();
fork();

-> this would result in 16 processes

parent
|-----
parent          1
|-----
parent    2          1          1.1
|-----  |-----  |-----  |-----
parent    2      2.1    1      1.2    1.1    1.1.1
|-----  |--  |--  |--  |--  |--  |--
```

danger: beware the "forkbomb"  
don't create infinitely many child processes

```
while (1) fork();    // never a good idea
```

exhausting the open process table can cripple the OS  
iLab has protection against users starting too many processes

How can we have multiple processes on a single-processor system?

- "time sharing"; some method for switching between running processes
  - cooperative multitasking
    - ("task" means "process" in this context)
    - each process runs for a bit and then yields control to the OS
    - the OS then resumes another process
    - problem: uncooperative processes can monopolize the CPU
  - preemptive multitasking
    - OS sets up a timer
    - each process runs for a short period of time
    - CPU interrupts program and returns control to OS
    - OS lets the next process have a slice of time

essential difference is who controls when we switch processes

- which is better?
  - cooperative multitasking is vulnerable to bad programs and bugs
    - one infinite loop can lock up the whole computer
  - preemptive multitasking requires more hardware support
    - preemption may occur at awkward times
    - e.g., a real-time system can't predict when it will be preempted

How does preemption occur?

-> Hardware interrupts or "traps"  
there are a bunch of related/similar ideas that have used different terms in different contexts, or used the same term in different ways

Basic idea: something happens where the CPU needs to respond to it immediately

- e.g., data from an IO device arrives
- run-time exception (division by zero, bad memory access)
- attempt to execute ill-formed/invalid instruction

If something happens, the CPU may interrupt the current process and transfer control to the OS

- current process state is saved
- control switches to OS code at a specified address (a "trap handler")
  - > trap handler will do something in response
  - copy data from IO device to a buffer in memory
  - terminate process with error condition
  - do nothing and resume process

For example, if we try to dereference a bad pointer (e.g., NULL)  
the CPU notices the attempt to read from an invalid address  
it switches to the trap handler for bad address errors  
the trap handler terminates our process with a SEGV signal

Typically, we (user programs) do not ever see traps or set trap handlers  
-> this is reserved for the OS

The OS can set alarms that will trap after some amount of time  
-> e.g., after 200 ms  
the trap handler can suspend the current process and have the scheduler resume another process

-> thus, preemptive multitasking

preemptive multitasking requires some hardware support, and is a bit more work than cooperative multitasking, but it is generally safer and more predictable

next time: signals

- signals are a way to interrupt a process
- processes can designate signal handlers that deal with signals when they arrive

start of multithreading