

Assignment 3 preview:

You will need to

- fork a child process
- use `execl` or `execv` to run another program
- hard mode:
  - redirect the program's standard output to a pipe
  - that is, the parent must read the output of the child

Simple way to send input from one fd to another

```
int infd, outfd, bytes;
char buf[SIZE];
```

```
while ((bytes = read(infd, buf, SIZE)) > 0) {
    write(outfd, buf, bytes);
}
```

----

recap:

- traps / hardware interrupts
  - a mechanism that allows the CPU to react to something immediately
  - suspends the current process and switches to some designated OS code
  - when the trap handler completes, OS can resume the suspended process
  - not something that user programs deal with directly

related idea: signals

- mechanism for communicating with a running process
- signals are sent to a process (from OS or other processes or the same process)
- they start out as "pending"
- normally, after some short period, they are "delivered" to the process

- for each signal, we declare a "disposition"
  - block the signal (leaves it pending; may get delivered later if signal is unblock)
  - ignore the signal
  - terminate process
  - execute a signal handler

a signal handler is a function that will be called when the signal is delivered

We can declare our own signal handlers using `signal`

```
#include <signal.h>

typedef void (*sighandler_t)(int);

sighandler_t signal(int signum, sighandler_t handler);
```

manual for `signal` function: `man 2 signal`  
manual for signals in general: `man 7 signal`

`signal` registers a signal handler

- There is a table somewhere that holds our process's disposition for each signal
- When the OS delivers a signal, it calls the appropriate signal handler

```
signal(int signal_number, sighandler_t signal_handler)

    sighandler_t is usually a void function that takes an int
or it could be:
    SIG_DFL - the default handler for this signal
    SIG_IGN - ignore this signal
```

When our process starts, it has a default disposition for every signal

- what that will do depends on the specific signal
  - some terminate the process
  - some terminate the process and create a core file
  - some stop the process (but it can be resumed later)
  - some are ignored

`man 7 signal` lists the default behaviors

To register a signal handler:

```
sighandler_t prev = signal(SIGINT, interrupt_handler);

prev is the previous signal handler, or SIG_ERR

-> any time you register a signal, you should check for SIG_ERR
```

To ignore `^C`

```
if (signal(SIGINT, SIG_IGN) == SIG_ERR) {
    // didn't work
}

// now we are immune to ^C
// not usually a good idea (makes it harder for user to stop the program)
// if we block SIGINT, then users will have to use SIGKILL to stop our program
//     SIGKILL is what kill -9 or kill -KILL sends
// we cannot set a handler for SIGKILL, so we can't do any cleanup if we
//     receive it
```

-> common technique is to intercept SIGINT, set a global variable, and then shut down cleanly

notes on signal handlers

-> signal handlers are functions that may be called from anywhere

- normally, the current function is interrupted and the signal handler is added to the call stack as though it had been called
- it is a function call that could happen at any time

-> it generally isn't a good idea to do a lot of work in a signal handler

- not all library functions are safe to call from a signal handler
  - e.g., you might get the signal in the middle of calling `printf` calling `printf` again may cause problems
- you may get another signal while you are executing the signal handler
- normally, a signal handler will not be interrupted by itself

`signal` does not behave consistently across Posix implementations

- the main difference has to do with what happens if you receive the same signal while the signal handler is running
  - possibility 1: use the default handler for that signal
  - possibility 2: block the signal until the handler completes
- portable code cannot assume which of these is being used
- GCC on the iLab uses #2

For portable code, use `sigaction`

- Posix (may not be available on non-Posix systems)
- more powerful and flexible
- more work to use

GCC Manual discussing signals  
[https://www.gnu.org/software/libc/manual/html\\_node/Signal-Handling.html](https://www.gnu.org/software/libc/manual/html_node/Signal-Handling.html)

Blocking signals

- in addition to signal handlers, our process has a "signal mask"
  - for each signal, the mask says that signal is blocked
  - a blocked signal stays pending; it does not get delivered
  - we can change the mask while we are running
  - if we unblock a signal, and there is a pending message for that signal, we receive it at that time

Waiting for signals

```
int pause(void);

pause suspends the current program until a signal is received
it returns the signal that was received, or -1 on error

e.g., if we didn't have sleep, we could implement our own using
alarm(some_time); // set an alarm -> receive SIGALRM when it is up
pause();          // wait until a signal is received

-> be sure to override the default behavior for SIGALRM first!
```

Termination signals

Different signals are used to terminate our process in different circumstances

`SIGHUP` - "hang up"; sent if the shell that started our process ends (e.g., we closed the window or logged out)

`SIGINT` - "interrupt from keyboard" - user typed `^C`

`SIGTERM` - "terminate", default signal for kill

`SIGQUIT` - "terminate and dump core"; used for debugging (type `^\`)

`SIGKILL` - "kill", terminate without cleanup

- processes cannot handle or ignore `SIGKILL`; it always terminates

Stop and continue

`SIGSTOP` - stop signal; cannot be handled or ignored

`SIGTSTP` - "typed stop"; sent when user types `^Z`; can be handled

`SIGCONT` - continue signal; sent when resuming a stopped process

Many other signals terminate by default

- usually because our program hit an error and can't safely proceed
  - division by zero
  - illegal memory access
  - malformed instruction
  - others (e.g., arithmetic errors)

- be careful when writing a handler for error condition
  - the error condition will still be there if the handler resumes
  - the only safe thing you can do is terminate the process or jump to somewhere else in the program (`siglongjmp`)

Aside: stopping and restarting processes in the shell

```
Use ^Z to stop the current process
Usually prints a message like: [1] Stopped  your_program
1 is the "job number"
    use jobs to get the current list of job
    processes started by the shell have job numbers
    all processes have PIDs
    Stopped is the current state of the process
```

```
To restart the process, use fg or bg
fg [job number]
    resumes process in the foreground
bg [job number]
    resumes process in the background (concurrent with shell)
```

Either give job number, or leave it out  
defaults to the most recent job

To start a process in the background, put `&` at the end of a command

```
$ ./long_process_to_finish &
[1] Running  ./long_process_to_finish
$
```

`kill` sends a signal to a process

```
kill PID    - send SIGTERM to process PID
kill -KILL PID - send SIGKILL to process PID
```

Use `%N` to get the PID of job `N`

```
$ ./long_process &
[1] 10234
$ kill %1
$
[1]+ Done    ./long_process
```

-----

```
#include <stdlib.h>
#include <unistd.h>
#include <stdio.h>
#include <signal.h>
```

```
volatile int signo = 0;
// marked volatile because it may change asynchronously
// that is, signals may be received at any time
```

```
// very simple signal handler
// note that it returns normally, so we can only use it with non-error signals
void handler(int signum)
{
    signo = signum;
}
```

```
// a very simple exit handler
void make_a_note(void)
{
    puts("We are in the exit handler");
}
```

```
int main(int argc, char **argv)
{
    // register an exit handler: make_a_note will be called after main returns
    atexit(make_a_note);

    // register some signal handlers
    // we can reuse handler because it will receive the signal
    signal(SIGHUP, handler);
    signal(SIGINT, handler);
    signal(SIGTERM, handler);
    signal(SIGCONT, handler);

    pause(); // stops the process until a signal is received
    // alternative:
    // while (signo == 0) { puts("Waiting"); sleep(1); }

    if (signo > 0) psignal(signo, "caught signal");

    return EXIT_SUCCESS;
}
```