

Command-line tricks

```
"file globs"

* -> matches any sequence of zero or more characters
If we use * in a command, the shell will look for files matching this pattern

*.txt -> replaced with the names of all files that end in .txt
foo*bar -> replaced with the names of all files that begin with foo and end with bar

? -> matches one character (wildcard)

    section??.txt -> matches section01.txt and section99.txt,
        but not section1.txt or section100.txt

By default, globs work in one directory (the working directory)
With /, we search in different directories

*.c    <- matches all C source files in the current directory
src/*.c <- all C source files in the subdirectory src
../*.c  <- all C source files in the parent directory

*/*.c   <- all C source files in any subdirectory
../*/*.c <- all C source files in any subdirectory of the parent directory

*.c
*/*.c
*/*/*.c

../src/module*/*.c
/usr/include/lib*/*.h
```

The shell replaces globs with lists of file names

If we type this
mv *.h include

The shell replaces it with
mv foo.h bar.h baz.h include

mv *.c *.bak <- sadly not possible
 <- mv never sees the globs, so it can't tell what we want to do

Shell scripts

```
Just bunch of shell commands in a (text) file
Mark as a executable
begin with #! followed by a path to the interpreter

Executing a shell script is just as if you had entered the commands yourself
```

```
---
#!/bin/bash
```

```
cp *.h backups
cp *.c backups
---
```

```
./mybackup.sh
```

Useful Unix commands

```
cat
- "concatenate"
- reads one or more files and prints to stdout

cat section*.txt
- combine all files whose names match the pattern and print
cat section*.txt > output
- combine all files and write to a file named "output"

more
- like cat, but pauses each time the screen fills up
- related program: less
  - does the same thing, but is fancier

- can give multiple arguments, or no arguments to read from stdin

    -> we can pipe the output of a program to more, and let more present it
        one screenful at a time

ps -ef          <- dumps information about all running processes
ps -ef | more   <- dumps the same information, one screen at a time

head [-number of lines] [files...]
- prints the first few lines of a file (or stdin)
- can optionally specify the number of lines (e.g. -20 to get 20 lines)

tail [-number of lines] [files...]
- prints the last few lines

Recall: ls -ltr lists all files in reverse chronological order
ls -ltr | tail      <- only prints the last few
ls -ltr | tail -1   <- only print the most recently modified file

file [files...]
- guess what kind of data a file contains
- uses a variety of methods, including looking for format markers
  and doing statistical analysis
- purely for your assistance; no programs depend on this output
- does not look at file names (extensions mean nothing)

wc [options] [files...]
- "word count"
- counts the number of characters, words, and lines in a file or files
- can request only certain counts (e.g., -l for just the number of lines)
- reads from stdin if no arguments

sort [files...]
- sorts lines alphabetically
  - concatenates all input files, and/or reads from stdin

uniq
- reads its input, but skips any line that is the same as the previous line
- read from multiple files and/or stdin
- use -c to print the number of times a line is repeated

grep [options] [pattern] [files...]

- multi-file text search using regular expressions
  - look for lines matching the pattern and print them
- can read from multiple files and/or stdin
  -> can use grep to search the output of a program
    ps -ef | grep lab
    ps -ef | grep -v root

regular expression syntax
* repeat zero or more times
+ repeat one or more times
. wildcard
[abcd] match any listed character
[a-z] match any listed character in the given range
^ match start of line
$ match end of line

[0-9]+ match any sequence of decimal digits

-> not the same syntax as file globs!

-v negates the grep (print the lines that do not match)
-c counts the number of lines that matched (per file)
-n prints file name and line number before the match

-> many options! so many!

cmp [file1] [file2]
says whether two files are the same
EXIT_SUCCESS and no output if the same
EXIT_FAILURE and some output if different

diff [file1] [file2]
compare two files and print all lines that are different

-> frequently used to compare versions of a source file

-> can use this to compare the output of a program to its expected output

./prog | diff expected -

- many options to control what is considered a difference and how it is reported
```

worth looking into: sed, awk

ps

- lists running processes
- use -e to get all processes
- many options to get more details, such as -f

top

- lists all running processes, sorted by CPU use
- updates the screen: live listing

Networks & inter-process communication

broadly: how do we send a message from one computer to another?

- more precisely, from one process to another

there are many ways for processes on the same computer to communicate

- > file system
- > pipes
- > use the OS to pass messages
- > shared memory

usually require both processes to be on the same computer
frequently, one process must be the parent of another

We want communication between processes

- running on different devices (no shared memory, disks, etc.)
- started independently of each other

Questions

- how is the communication organized
 - send individual messages?
 - stream of bytes?
- how are messages transferred from one process to the other?
- how do the processes identify each other?

Any networking system must be able to answer these questions

- there have been many networking systems that have had different answers

Typical designs are layered

- different subsystems have different areas of responsibility
- standard interfaces used to access these systems
 - > sockets are one such interface

socket interface is very general

- > we can use the same interface for many different kinds of network
- > a typical program will only use a small number of networks (e.g., one)