

In POSIX C, we have two families of IO operations

Posix operations / syscalls / non-buffered IO

```
read() and write() (also send() and recv())

- low-level calls
- provide the same interface for files, sockets, pipes, etc
- precise control
  - the data I send using write() is sent to the OS / no longer my responsibility
  - when I receive data using read(), I can set a specific maximum

-> no built-in formatted IO
    printing a decimal integer requires allocating a local buffer
    and either writing a function or using sprintf()
```

C operations / buffered IO

```
fread(), fwrite(), getc(), putc()
fprintf(), fscanf()

these maintain a local buffer
  data written using putc() and fprintf() may not get sent immediately
  to the OS

  when we call getc() or fscanf(), the library may prefetch data and
  store it in the buffer

-> less control about when data is sent and received
-> potentially fewer system calls (= better performance)
  -> calling getc() in a loop is far more efficient than calling
      read(fd, &ch, 1)

If you want to use buffered IO with a socket, you will probably want
separate read and write buffers

FILE *fp = fdopen(socket, "r")
// create a FILE for a socket in read/write mode
// only has one buffer, shared by reading and writing operations
// after we write, we must fflush() before reading
// after we read, must reset the file pointer before writing

Instead, when using buffered IO on a socket, duplicate the socket and
create two separate FILES

int sock2 = dup(sock);
// should confirm sock2 != -1
FILE *fin = fdopen(sock, "r");
FILE *fout = fdopen(sock2, "w");
// should confirm fin != NULL and fout != NULL

when you call fclose(), it calls close() on the file descriptor

    fclose(fin);
    fclose(fout);

fscanf(fin, "%d" &len);
// reads bytes until the first non-digit

Note: sockets do not have all the features of files
-> we can't skip around; no way to skip forward or back
-> we don't necessarily know how many bytes we have written

recall: writing using buffered IO does not necessarily send data to the OS
immediately

    fflush(fout);    // sends anything in the buffer to the OS

fprintf(fout, "OKG\n%d\n%s\n", strlen(value)+1, value);
// puts data into the buffer
fflush(fout);
// sends buffer contents to OS -> to client
```

Sockets do not behave exactly like files

```
read() with files usually gives us all the data we want
read() with sockets gives us all the data currently available
-> may be less than we asked for
-> we have no control over how quickly data arrives from the other
    party
-> if the connection is bad, we could get data 1-byte at a time!
```

write() with files pretty much always writes all the bytes given

write() with sockets usually does, but might not

```
-> other party may have closed the connection
-> we could have been interrupted by a signal
```

-> for maximum safety, any time we call write() we need to check how
much was actually written, and possibly call write() again to rewrite
the remainder

```
char *buf;
int buflen, bytes, written = 0;

while (written < buflen) {
    bytes = write(fd, buf + written, buflen - written);
    if (bytes < 1) { some sort of error handling }

    written += bytes;
}

-> we don't usually need to do this, because write() usually writes
    everything
-> we should still check the return value in case we were interrupted
    by a signal, or the connection closed/file became unwritable
```

Fuzzy thinking about sockets

- * our protocol is described in terms of messages
 - clients sends a request
 - server sends a response
- * the TCP model gives us two streams of bytes
 - TCP does not guarantee that we will get a whole message at once
 - TCP does not guarantee a break between messages

When we read, we don't want to read too much

- if the client sent another request without waiting for our response,
we could read part of the second message
 - > now we have to hold onto it until we finish dealing with the first
one
- if we ask for more bytes than the client sent, but the client is waiting
for our response before it sends any more, then we deadlock
 - server is waiting for the client
 - client is waiting for server
- buggy client sends message that is too short
- buggy server asks for more bytes than it should
- we defend against server bugs by writing good code
- we defend against short messages from the client by looking for the
terminating newline

```
GET\n5\n\day\n

terminating \n comes after 4 characters, not 5
server should immediately detect this, and not try to read further
-> send ERR\nLEN\n and close connection

GET\n3\day\n

server should not read past the 'y'
-> send ERR\nLEN\n and close connection

GET\n4\nday

no terminating \n
-> maybe not even a problem; client could still send the \n
    later
-> defending against this requires more advanced features
-> we won't consider this scenario in this class
-> you can't fix this with the tools I have given you,
    so don't try
```

Note about netcat (nc)

- we can use nc to pretend to be a client for Project III
- nc is one possible interpretation of the protocol
 - nc may not interleave input and output precisely
 - if data arrives while nc is waiting for input, it may not show us
the data immediately
- if you want to know the exact sequence of events, write a client

Multithreading in servers

echos.c, and your Project III, start a thread for each connection request

```
loop {

    struct arg_t args = malloc(sizeof(struct arg_t));

    connection_fd = accept(listening_fd, NULL, NULL);
    // block until a remote host tries to connect
    if (connection_fd < 0) ...

    args->fd = connection_fd;

    err = pthread_create(&worker_id, NULL, worker_function, args);
    // worker_fun will eventually close connection_fd and free args
    if (err) ...

    pthread_detach(worker_id);
    // thread will clean up after itself; no return value
    // the main thread doesn't need to remember how many threads are
    // running, and does not need to join them
}
```

As written, the only way to stop this server is to terminate it with
SIGINT, SIGTERM, SIGKILL, etc

-> server shuts down immediately, closes all sockets, terminates all
threads

How can we make this cleaner?

- we can use a signal handler to catch SIGINT and/or SIGKILL, etc.
- many blocking system calls will return early if you receive a signal
while you are blocked
- they may also be set to auto-resume by default

One possible strategy

- install a signal handler
- have some flag that indicates whether the signal has been received
- accept() will return -1 if it is interrupted by a signal
- have the main loop exit if the signal was received

```
while (running) {
    fd = accept(listener_fd, &addr, &addrlen);
    if (fd < 0) continue;

    .. start up thread to handle connection
}

// we don't get here until after the signal arrives
```

A few points to consider

- when we return from main(), the whole process stops and all threads are
terminated
- if we call pthread_exit(), only that thread ends
 - if the main thread exits, the process won't terminate until every
thread is finished
- any exit handlers will get called after the last thread finishes
e.g.,
atexit(cleanup_data);
- signals are sent to any thread that is not blocking the signal
 - each thread has its own signal mask (= which signals are blocked)
 - all threads have the same signal dispositions (SIG_IGN, SIG_DEF, functions)
 - accept() will only break out of the block if the main thread gets the
signal
- we need to arrange things so that only the main thread receives the
signals of interest
- when a thread is created, it inherits its parent's signal mask
- so:
 - block SIGINT
 - spawn child thread
 - unblock SIGINT