

# CS214-system programming

- Section 03/08 recitation 04

Yunhe Gao

yg397@scarletmail.rutgers.edu

# Content

- File operation
- Directory

# File types in Linux system

- A file is a sequence of bytes
- All I/O devices (network, disks, ...) are modeled as files to generalize all input and output operations as uniform ones
- Regular file: .txt, .out, .o, .c, ...
- Directory: a file composing a set of file links
- Another directory can be an element in the set recursively

## "Everything is a File" and Types of Files in Linux

Normal	-	Normal file
Directories	d	Normal directory
Hard link	-	additional name for existing file
Symbolic link	l	Shortcut to a file or directory
Socket	s	Pass data between 2 process
Named pipe	p	like sockets, user can't work directly with it
Character device	c	Processes character hw communication
Block device	b	Major and minor numbers for controlling dev

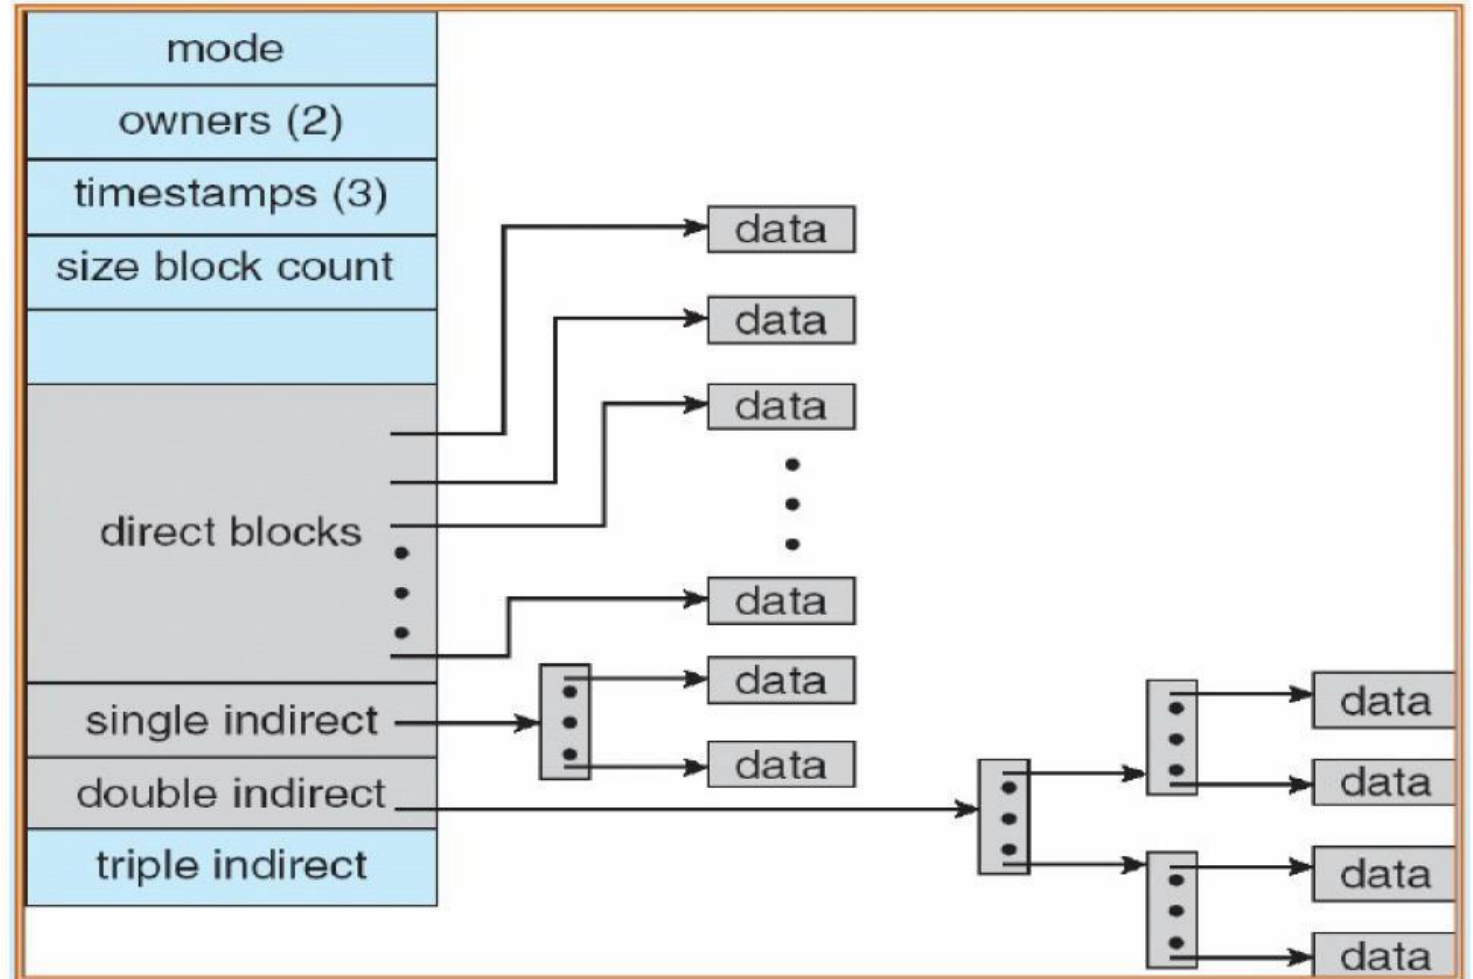
# inode

Each file on the disk has an inode associated with it.


An inode is a data structure on a filesystem that stores meta-information about a file or directory.

Consists of:

- File metadata (mode, owner, info, size, etc.)
- Pointers to data blocks
  - Direct mapped pointers to data blocks
  - Indirect pointers
    - Point to blocks that contain pointers



inode



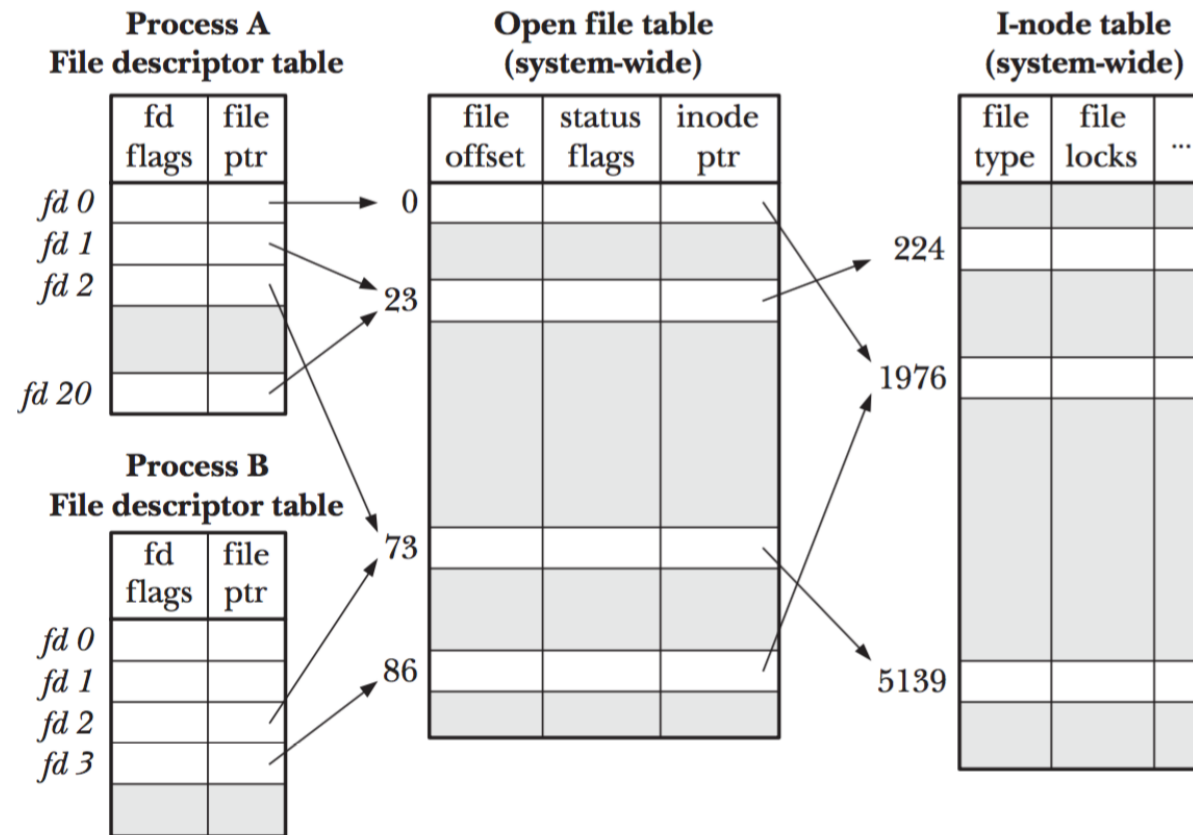
# File operations in Linux

- Three basic types of operations on a file
  - Read (r)
  - Write (w)
  - Execute (x)
- A permission of an operation is associated with the relationship between current user and the file
- Three types of relationships
  - Owner (u)
  - Users in the group where the owner is in (g)
  - Users in other groups (o)

# File operations

- File descriptor:
  - Unix / Linux I/O functions: access I/O devices via system calls
  - File descriptor is used for file operations, no stream is needed
  - Open(), close(), read(), write(), ...
  - Low-level operations are performed directly using file descriptor
- FILE pointer:
  - Standard I/O functions: access I/O devices in a higher-level
  - A stream (represented with a FILE object) is needed to associate with an opened file
  - Implemented by invoking Linux / Unix I/O functions
  - printf(), scanf(), fopen(), fread(), fscanf(), fgets(), ...
  - Streams interface could provide powerful formatted input and output functions

# An opened file in Linux



**Figure 5-2:** Relationship between file descriptors, open file descriptions, and i-nodes



# Open

- Open a file: kernel does the work for the program
- After opening, a file descriptor is returned to the program
- `int open(char *path, int flags, mode_t mode);`
- Returns: new file descriptor if OK, -1 otherwise
- Import library: `fcntl.h`

# Close Function and Errno

- Close an opened file represented by a file descriptor
- Prototype:  

```
int close(int fd);
```
- Import library: `unistd.h`
- Returns: 0 on success, -1 on error, ***errno*** is set appropriately
- Errno: indicate what goes wrong in the event of an error in the system
- Errno is set by system calls and some library functions

# Read Function

- Read: copies  $\leq n$  bytes from the current position of *fd* to memory location *buf*
- Import library: unistd.h
- Prototype:

```
ssize_t read(int fd, void *buf, size_t n);
```

- Return number of bytes read if OK, 0 on EOF, -1 on error
- `ssize_t`: signed long in essence, which could be a negative number (-1 on error)
- `size_t`: unsigned long in essence, which should satisfy that is bigger or equal to 0

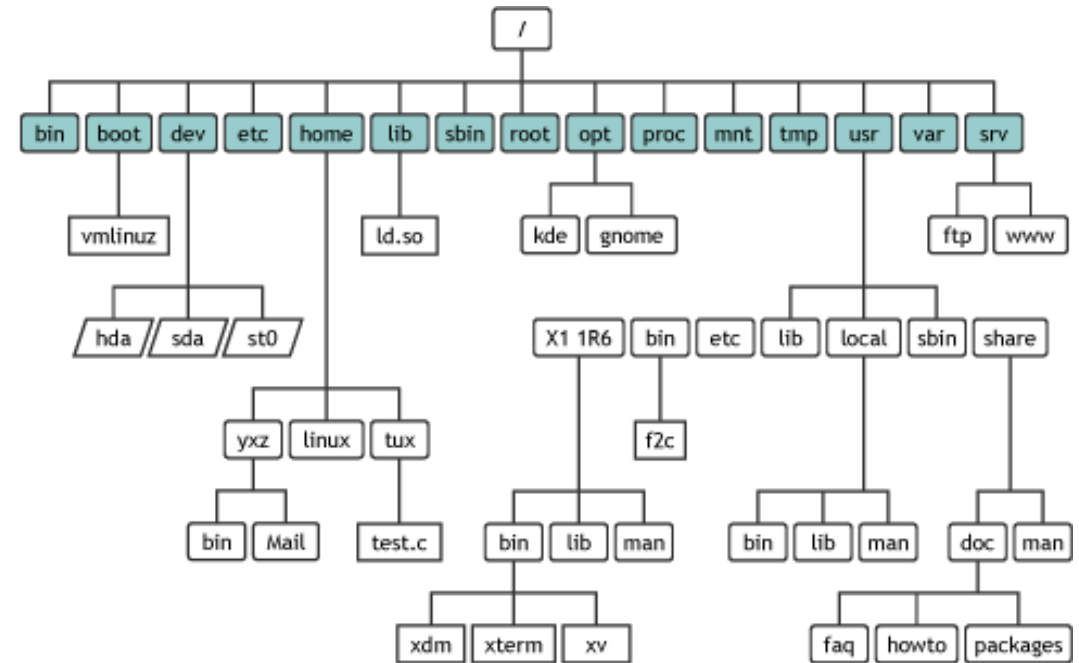
# Write Function

- Write: copies  $\leq n$  bytes from the memory location *buf* to the current file position of *fd*
- Import library: unistd.h
- Prototype of writes

```
ssize_t write(int fd, const void *buf, size_t n);
```
- Returns number of bytes written if OK, -1 on error

# Directory

- Directory: special file that contains directory entries
- In Linux, directories are organized in form of a tree and there exists a root directory
- 3 permissions associated with a directory
- Read: get the list of file links
- Write: add/remove a file link
- Execute: enter into the directory (*cd* command), access files inside it and view the metadata of the files (*/s -/* command)



<http://researchhubs.com/post/computing/linux-cmd/linux-directory.html>

# Directory

- It is a special file that contains directory entries
- How to find if a file is a regular file or a directory, and other information?
  - Use `stat`

```
int stat(const char *pathname, struct stat *statbuf);
```

# stat

- `int stat(const char *pathname, struct stat *statbuf);`

```
The stat structure
All of these system calls return a stat structure, which contains the following fields:

    struct stat {
        dev_t    st_dev;        /* ID of device containing file */
        ino_t     st_ino;       /* Inode number */
        mode_t    st_mode;      /* File type and mode */
        nlink_t   st_nlink;     /* Number of hard links */
        uid_t     st_uid;       /* User ID of owner */
        gid_t     st_gid;       /* Group ID of owner */
        dev_t     st_rdev;      /* Device ID (if special file) */
        off_t     st_size;      /* Total size, in bytes */
        blksize_t st_blksize;    /* Block size for filesystem I/O */
        blkcnt_t  st_blocks;    /* Number of 512B blocks allocated */

        /* Since Linux 2.6, the kernel supports nanosecond
           precision for the following timestamp fields.
           For the details before Linux 2.6, see NOTES. */

        struct timespec st_atim; /* Time of last access */
        struct timespec st_mtim; /* Time of last modification */
        struct timespec st_ctim; /* Time of last status change */

        #define st_atime st_atim.tv_sec      /* Backward compatibility */
        #define st_mtime st_mtim.tv_sec
        #define st_ctime st_ctim.tv_sec
    };
```

# stat

The following mask values are defined for the file type:

S_IFMT	0170000	bit mask for the file type bit field
S_IFSOCK	0140000	socket
S_IFLNK	0120000	symbolic link
S_IFREG	0100000	regular file
S_IFBLK	0060000	block device
S_IFDIR	0040000	directory
S_IFCHR	0020000	character device
S_IFIFO	0010000	FIFO

Thus, to test for a regular file (for example), one could write:

```
stat(pathname, &sb);
if ((sb.st_mode & S_IFMT) == S_IFREG) {
    /* Handle regular file */
}
```

Because tests of the above form are common, additional macros are defined by POSIX to allow the test of the file type in st\_mode to be written more concisely:

```
S_ISREG(m)  is it a regular file?
S_ISDIR(m)  directory?
S_ISCHR(m)  character device?
S_ISBLK(m)  block device?
S_ISFIFO(m) FIFO (named pipe)?
S_ISLNK(m)  symbolic link? (Not in POSIX.1-1996.)
S_ISSOCK(m) socket? (Not in POSIX.1-1996.)
```

The preceding code snippet could thus be rewritten as:

```
stat(pathname, &sb);
if (S_ISREG(sb.st_mode)) {
    /* Handle regular file */
}
```



# Directory Handling

- `DIR *opendir(char *path);`
  - Returns a pointer to an abstract DIR structure
  - Returns NULL and sets errno on failure
- `int closedir(DIR *directory_pointer);`
  - Close the directory & deallocate struct
  - Return -1 and sets errno on failure
- `struct dirent *readdir(DIR *directory_pointer);`
  - Obtain the next entry from the directory
  - Return NULL if there are no further entries

# Directory Handling

## SYNOPSIS

```
#include <sys/types.h>
#include <dirent.h>
```

```
DIR *opendir(const char *name);
struct dirent *readdir(DIR *dirp);
```

```
struct dirent {
    ino_t      d_ino;        /* inode number */
    off_t      d_off;        /* not an offset; see NOTES */
    unsigned short d_reclen; /* length of this record */
    unsigned char d_type;    /* type of file; not supported
                             by all filesystem types */
    char        d_name[256]; /* filename */
};
```

DT_BLK	This is a block device.
DT_CHR	This is a character device.
DT_DIR	This is a directory.
DT_FIFO	This is a named pipe (FIFO).
DT_LNK	This is a symbolic link.
DT_REG	This is a regular file.
DT SOCK	This is a UNIX domain socket.
DT_UNKNOWN	The file type is unknown.

```
1 #include <stdlib.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <dirent.h>
5 #include <string.h>
6
7 int main()
8 {
9     DIR *dirp = opendir(".");
10
11     struct dirent *dp;
12
13     while ((dp = readdir(dirp))) {
14         printf("Got file %llu: %s\n", (unsigned long long) dp->d_ino, dp->d_name);
15     }
16
17     closedir(dirp);
18
19     return EXIT_SUCCESS;
20 }
21
```

# Thanks

- [yg397@scarletmail.rutgers.edu](mailto:yg397@scarletmail.rutgers.edu)